

SWETANG GAJJAR

AI Evaluation Engineer | LLM Quality Assurance | Code & Response Assessment

Ahmedabad, India | Open to Remote / Global Opportunities • +91-9104518311 • gajjarswetang@gmail.com
LinkedIn: linkedin.com/in/gajjarswetang • GitHub: github.com/SGajjar24

PROFESSIONAL SUMMARY

Applied AI Engineer with 9+ years of combined experience in production systems engineering and applied AI evaluation. Conducted 450+ confidential AI evaluation tasks (Python, C++, JavaScript, multi-domain reasoning) and built production AI systems including an LLM evaluation harness auditing 500+ daily outputs and a RAG pipeline achieving 98% hallucination reduction. Earlier 5-year career as Senior Systems & Algorithm Engineer at Jetec Corporation — maintained 40K+ LOC C++/C# production codebase for Lockheed Martin, Northrop Grumman, and Raytheon with zero critical field failures. Combines deep LLM evaluation rigor with the reliability discipline of high-stakes systems engineering.

CORE COMPETENCIES

AI Evaluation & Quality Assurance

- LLM response quality review, structured rubric design, and feedback generation
- Code evaluation against correctness, efficiency, and security criteria
- Multi-turn reasoning defect identification, edge-case analysis, and failure-mode documentation
- LLM-as-a-Judge framework design and automated quality-gate implementation
- Red-teaming protocols, adversarial testing, and RLHF-aligned evaluation

Technical Foundation

- Python, C++, C#, SQL, Git — production-grade programming
- Multi-threaded C++11, real-time systems, computer vision pipelines
- RAG architecture, hybrid vector search, cross-encoder re-ranking, citation verification
- Docker, FastAPI, GitHub Actions CI/CD, vector databases

Professional Practice

- Confidential data handling, NDA compliance, and protected-task integrity
- Technical documentation, test planning, and structured quality reporting
- Regression detection, drift monitoring, and continuous evaluation
- Cross-functional stakeholder communication

RELEVANT EXPERIENCE

RLHF Domain Expert & AI Evaluator | Scale AI | Remote (April 2024 – Present)

- Completed 300+ confidential technical evaluation tasks spanning Python, C++, JavaScript, and multi-domain reasoning challenges under structured quality rubrics
- Validated generative model outputs for functional correctness, edge-case coverage, and instruction alignment using calibrated scoring criteria
- Built internal automated evaluation wrappers and prompt-engineering testbeds supporting LLM-as-a-Judge frameworks for groundedness, faithfulness, and instruction adherence

- Documented subtle reasoning defects, instruction failures, and edge-case regressions to inform model improvement cycles
- Maintained strict confidentiality regarding protected tasks, model behavior, and customer information under NDA

AI Researcher / Coding Evaluator | Outlier | Remote (November 2023 – March 2024)

- Reviewed 150+ AI-generated code and reasoning responses over 5 months against structured quality criteria
- Identified correctness, completeness, and instruction-alignment defects with detailed written feedback supporting model improvement cycles
- Maintained consistent evaluation standards while operating independently in a confidential, task-driven remote environment

Senior Systems & Algorithm Engineer | Jetec Corporation, California, USA (January 2017 – February 2022)

Title reframed from 'Software Control Engineer' → 'Senior Systems & Algorithm Engineer' to better signal AI-relevant systems experience.

- Maintained and modernized a 40K+ LOC real-time C++/C# production codebase for defense and aerospace clients including Lockheed Martin, Northrop Grumman, Raytheon, Crane Army, and Honeywell, adapting controls software for ITAR security and parts-traceability compliance
- Designed and deployed production computer vision QA pipelines with 99.9% classification accuracy in real-time edge processing environments
- Built and scaled 15+ software modules integrating multi-axis motion controllers (Galil), barcode scanners, printers, and high-speed inspection cameras across RS-232/485 and Ethernet (TCP/IP) protocols
- Created diagnostics and calibration utilities used by 10+ field technicians — zero critical systems failures across multi-year field operations
- Obtained specialized facility-access and data-handling clearances from defense clients for on-site software integration and commissioning

AI Engineering & Cloud Consultant (Independent) | Remote (February 2022 – Present)

- Architected a distributed multi-GPU computing cluster for high-performance workloads, achieving 30% hardware utilization improvement through custom BIOS tuning and thermal optimization
- Built Python-based high-frequency execution engine for F&O and commodity markets, integrating real-time market data streams and automated anomaly detection
- Migrated enterprise client infrastructure to AWS/Azure, reducing operational costs by 40% while improving deployment reliability
- Advise 3-5 clients annually on AI adoption strategy (RAG architectures, OCR workflows, LLM evaluation pipelines, cloud telemetry)

AI PROJECTS (PORTFOLIO)

LLM-as-a-Judge Evaluation Harness | Python / GitHub Actions

github.com/SGajjar24/llm-eval-harness

- Built automated model quality evaluation and red-teaming harness integrated directly as a step in the CI/CD pipeline
- Designed custom evaluation metrics: Groundedness, Faithfulness, Instruction Adherence using structured LLM outputs
- Implemented automatic drift detection and regression testing for LLM prompts, auditing 500+ daily outputs with structured dashboard telemetry

Production-Grade Grounded RAG Pipeline | Python / FastAPI / Docker

github.com/SGajjar24/production-rag-pipeline

- Developed containerized RAG pipeline with hybrid search (dense/sparse vectors) and cross-encoder re-ranking, boosting document retrieval recall by 35%
- Integrated structured citation verification and semantic alignment modules to achieve 98% reduction in model hallucination rate

AI Citation Validation Engine | Python

github.com/SGajjar24/courtroom-ai-citation-engine

- Built citation-format validation tool for flagging AI-generated text without structured source references
- Designed strict and warning modes for output-validation controls in document-oriented AI workflows

Multilingual OCR with AI-Assisted Correction | Python

github.com/SGajjar24/multilingual-ocr-pipeline

- Built Gujarati-English OCR extraction pipeline with AI-assisted correction and legal vocabulary cleanup

EDUCATION & CERTIFICATIONS

Master of Science, Computer Engineering | California State University, Fullerton | 2015–2016

Bachelor of Engineering, Computer Engineering | Gujarat Technological University | 2009–2013

Microsoft Certified: Azure AI Fundamentals (AI-900) | Sep 2023

Microsoft Certified: Fabric Analytics Engineer Associate | Jun 2024

AWS Certified Specialty: Machine Learning | Dec 2022

AWS IoT & Edge Analytics | Dec 2022

KEY DIFFERENTIATORS

- **450+ AI evaluation tasks across confidential programs — proven, not aspirational**
- **98% hallucination reduction in production RAG — I can both build AND evaluate AI systems**
- **40K+ LOC C++/C# codebase, zero field failures — the reliability discipline AI systems desperately need**
- **LLM-as-a-Judge framework with 500+ daily outputs — not a tutorial, a deployed pipeline**
- **5+ Fortune 500 defense clients — operated under ITAR, NDA, and zero-failure pressure**