Open to global opportunities · US · Canada · EU · UK · AU

Building AI systems that have to *work*.

I'm **Swetang Gajjar** — an Applied AI Engineer with 9+ years of production engineering experience, now focused on **LLM evaluation**, **RAG architecture**, and shipping AI systems that don't break.

See my work GitHub LinkedIn

0
Years building production systems
0
Confidential AI evaluation tasks
0
Hallucination reduction in production RAG
0
LOC production C++/C# codebase
SCROLL

AI systems don't fail at the model. They fail at the seam between the model and the world. **That's where I work.**

A WORKING PRINCIPLE

/01 — WHAT I DO

Bridging systems rigor with modern AI

Two and a half years of evaluation work + 5 years of production engineering + 4 open-source AI projects. I build, evaluate, and ship.

I evaluate generative AI systems and ship the infrastructure around them.

Over the past 2+ years, I've evaluated **450+ AI-generated technical responses** across Python, C++, JavaScript, and multi-domain reasoning under confidential programs at Scale AI and Outlier. I've also built and shipped production AI systems — including a containerized RAG pipeline that cut hallucination rates by 98% and a CI/CD-integrated evaluation harness auditing 500+ daily LLM outputs.

Before moving into AI, I spent 5+ years at **Jetec Corporation** in California maintaining a **40K+ LOC real-time C++/C# production codebase** for defense and aerospace clients (Lockheed Martin, Northrop Grumman, Raytheon, Crane Army, Honeywell) — where I developed the debugging discipline and zero-failure engineering mindset that now shapes how I evaluate and ship AI systems.

M.S. Computer Engineering from **Cal State Fullerton**. Currently in Ahmedabad, India. Open to global remote + relocation (US, Canada, EU, UK, Australia).

01

LLM Evaluation & Quality

Multi-turn reasoning, code, and safety evaluation. Structured rubrics, LLM-as-a-Judge frameworks, red-teaming, RLHF-aligned feedback.

EVALUATION

02

Applied AI & RAG

Containerized RAG with hybrid dense/sparse search, cross-encoder re-ranking, structured citation verification, AI-assisted OCR pipelines.

RAG · OCR

03

Production Engineering

Multi-threaded C++ systems, real-time computer vision, 15+ production modules, zero critical field failures across 5+ years.

C++ · SYSTEMS

04

AI Infrastructure

Docker, FastAPI, GitHub Actions CI/CD, multi-GPU cluster tuning, AWS/Azure migration (40% cost reduction for clients).

MLOPS · CLOUD
/02 — RIGHT NOW

Currently building

Live status of what I'm working on, what I'm exploring, and what's next on the roadmap.

ACTIVE

Multi-agent eval harness

Extending llm-eval-harness to support multi-agent system evaluation — looking at how agent trace data can be scored for correctness and reliability.

PythonTracingLLM-as-Judge
ACTIVE

Open-source RAG improvements

Production-rag-pipeline v2 — adding cross-encoder re-ranking, better citation UX, and post-deployment drift monitoring.

FastAPIVector DBDocker
PENDING

Job search — AI Eval / Applied AI

Actively interviewing for AI Evaluation, Applied AI, Forward Deployed AI, and LLM Ops roles. Open to remote + global relocation.

Open to offersRemoteRelocation
/03 — WHERE I'VE WORKED

A 9-year arc from industrial systems to AI

Apr 2024 — Present **CURRENT**

RLHF Domain Expert & AI Evaluator

Scale AI · Remote

- Completed 300+ confidential technical evaluation tasks spanning Python, C++, JavaScript, and multi-domain reasoning challenges under structured quality rubrics.

- Validated generative model outputs for functional correctness, edge-case coverage, and instruction alignment using calibrated scoring criteria.
- Built internal automated evaluation wrappers and prompt-engineering testbeds supporting LLM-as-a-Judge frameworks for groundedness, faithfulness, and instruction adherence.
- Documented subtle reasoning defects, instruction failures, and edge-case regressions to inform model improvement cycles.

Feb 2022 — Present Independent

AI Engineering & Cloud Consultant

Self-employed · Remote

- Architected a distributed multi-GPU computing cluster for high-performance workloads, achieving 30% hardware utilization improvement through custom BIOS tuning and thermal optimization.
- Built Python-based high-frequency execution engine for F&O and commodity markets, integrating real-time market data streams and automated anomaly detection.
- Migrated enterprise client infrastructure to AWS/Azure, reducing operational costs by 40%.
- Advise 3-5 clients annually on AI adoption strategy (RAG architectures, OCR workflows, LLM evaluation pipelines, cloud telemetry).

Nov 2023 — Mar 2024

AI Researcher / Coding Evaluator

Outlier · Remote

- Reviewed 150+ AI-generated code and reasoning responses over 5 months against structured quality criteria.
- Identified correctness, completeness, and instruction-alignment defects with detailed written feedback supporting model improvement cycles.
- Maintained consistent evaluation standards while operating independently in a confidential, task-driven remote environment.

Jan 2017 — Feb 2022 5+ years Fortune 500 defense

Senior Systems & Algorithm Engineer

Jetec Corporation · Tustin, California, USA

- Maintained and modernized a **40K+ LOC real-time C++/C# production codebase** for defense and aerospace clients including Lockheed Martin, Northrop Grumman, Raytheon, Crane Army, and Honeywell.
- Designed and deployed production computer vision QA pipelines with **99.9% classification accuracy** in real-time edge processing.
- Built and scaled 15+ software modules integrating multi-axis motion controllers, barcode scanners, and high-speed inspection cameras.

- Created diagnostics used by 10+ field technicians — **zero critical systems failures** across multi-year field operations.

/04 — WHAT I'VE SHIPPED

Open-source AI systems in production

Every project is open-source on GitHub, deployed somewhere, and demonstrates a different facet of the AI evaluation + applied AI stack.

PRIMARY · LLM EVALUATION

LLM-as-a-Judge Evaluation Harness

Automated model quality evaluation pipeline with custom **Groundedness, Faithfulness, and Instruction Adherence** metrics. Integrated as a CI/CD step with structured dashboard telemetry auditing **500+ daily LLM outputs**.

500+
DAILY LLM OUTPUTS
3
CUSTOM METRICS
CI/CD
GITHUB ACTIONS

Python GitHub Actions OpenAI API Structured Outputs

View repository

◆

RAG

Production-Grade Grounded RAG

Containerized RAG with hybrid dense/sparse vector search, cross-encoder re-ranking, and structured citation verification.

35%
RECALL BOOST
98%
HALLUCINATION CUT

Docker
CONTAINERIZED

Python FastAPI Docker Vector DB

View repository

QUALITY

AI Citation Validation Engine

Citation-format validation tool for flagging AI-generated text without structured source references. Strict + warning modes for document AI workflows.

2
MODES
100%
LOCAL-FIRST
CLI
+ LIBRARY

Python Regex Document AI

View repository

OCR

Multilingual OCR Pipeline

Gujarati-English OCR extraction with AI-assisted correction and legal vocabulary cleanup. Real-world document AI for low-resource languages.

2
LANGUAGES
AI
POST-PROCESS
Legal
DOMAIN-TUNED

Python Tesseract NLP

View repository

/05 — WORDS I THINK ABOUT

Three things I keep coming back to

I.

Reliability is a feature.

The most underrated capability in production AI isn't clever prompting or a bigger model. It's the system that runs reliably at 3am when nobody's watching. That's the bar.

II.

Evaluation is infrastructure.

Most teams treat evals as a one-off QA task. The best teams treat them as code: versioned, automated, in CI/CD, and as a first-class signal for what to build next.

III.

Production teaches what demos hide.

Five years of shipping systems that can't fail taught me more about AI quality than any benchmark. The gaps only show up in production.

/06 — GET IN TOUCH

Let's build reliable AI together.

Open to **Applied AI Engineer, AI Evaluation Engineer, Forward Deployed AI Engineer, LLM Ops/MLOps, AI Product Engineer**, and **Technical Program Manager (AI)** roles. Remote, hybrid, or on-site — including US, Canada, EU, UK, and Australia.

EMAIL
gajjarswetang@gmail.com

LINKEDIN
linkedin.com/in/gajjarswetang

GITHUB

github.com/SGajjar24

Based in Ahmedabad, India · Open to global opportunities · Prior US H-1B cap count preserved · M.S. Cal State Fullerton

Back to top ↑